

Reinforcement Learning

Dynamic Programming

Prof. James Brusey

June 4, 2023

Why study Dynamic Programming?

- ▶ DP assumes perfect model and is computationally expensive
- ▶ Still important theoretically, though
- ▶ Basis for understanding

Expectation (reminder)

- ▶ Expected value is the average outcome of a random event given a large number of trials
- ▶ Where all possible values are equally likely, this is simply the mean of possible values
- ▶ Where each x_i occurs with probability $p(x_i)$, we have the sum

$$\mathbb{E}[x] = \sum_i p(x_i)x_i$$

- ▶ Quick quiz: what's the expected value of a fair, 6-sized dice?

Expectation with a subscript?

- ▶ When we have a subscript, that means according to that probability measure

$$\mathbb{E}_y[x] = \sum_i p_y(x_i)x_i$$

where p_y is a probability measure for x that might differ from the sampling probability.

Bellman optimality equation

- ▶ Dynamic programming is based on solving Bellman's optimality equation

$$v_*(s) = \max_a \mathbb{E} [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a]$$

- ▶ why is this equation so important and what do the variables mean?

Policy evaluation (prediction)

Evaluation asks:

- ▶ Given a policy π , what is the expected value v_π of each state s ?

$$v_\pi(s) \doteq \mathbb{E}_\pi[G_t | S_t = s]$$

Policy evaluation (2)

- We can substitute for the G_t based on

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$$

or more simply

$$G_t = R_{t+1} + \gamma G_{t+1}$$

Policy evaluation (3)

We should get

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$


where

- ▶ $p(s', r|s, a)$ is the transition probability
- ▶ $\pi(a|s)$ is the probability of taking action a when in state s
- ▶ r is the immediate reward
- ▶ γ is the discount factor

Policy evaluation (4)

For this simple recycle robot, write down the set of equations for $v_{\pi}(s)$ assuming that π is a uniform distribution over actions.

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-



Policy evaluation (5)

Here is the equation for $v_\pi(\text{high})$:

$$\begin{aligned} v_\pi(\text{high}) &= \frac{1}{2}(r_{\text{wait}} + \gamma v_\pi(\text{high})) + \\ &\quad \frac{1}{2}\left(\alpha(r_{\text{search}} + \gamma v_\pi(\text{high})) + \right. \\ &\quad \left.(1 - \alpha)(r_{\text{search}} + \gamma v_\pi(\text{low}))\right) \end{aligned}$$

LP solution

Note that we have equations in the form:

$$v(a) = k_1 v(a) + k_2 v(b) + k_3$$

which can be converted into

$$-k_3 = (k_1 - 1)v(a) + k_2 v(b)$$

or

$$\begin{pmatrix} -k_3 \\ \vdots \end{pmatrix} = \begin{bmatrix} (k_1 - 1) & k_2 \\ \vdots & \vdots \end{bmatrix} v$$

Which can be solved by inverting the matrix. *Do try this at home!*


Iterative policy evaluation

v_k for the random policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0


greedy policy w.r.t. v_k



random policy


$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



$k = 2$


0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0



Iterative policy evaluation


$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0



$k = 10$


0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



optimal
policy

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



Approximate methods

- ▶ It is also possible to find the solution *iteratively* by starting off with some guesses for v and updating according to the equations.
- ▶ This approach forms the basis of many other methods and for large MDPs we will never attempt using LP

Policy improvement

- ▶ Given that we can evaluate any policy, we can therefore *improve* it by updating the policy so that it takes the action that maximises the resulting value
- ▶ We define the expected value of taking an action a and then following π from then on

$$q_\pi(s, a) \doteq \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a]$$

- ▶ which expands to

$$= \sum_{s',r} p(s', r | s, a)[r + \gamma v_\pi(s')]$$

Policy improvement (2)

- ▶ If we can find some a for which $q_{\pi}(s, a) > v_{\pi}(s)$, then we can *improve* π by adjusting it to use a
- ▶ If we cannot find any improvement for all states then π must be optimal
- ▶ Discuss:
 - ▶ what aspects of MDPs are relied on to ensure this is true?

Policy iteration algorithm

S1: Initialisation

S2: Policy evaluation

▶ Loop:

- ▶ $\Delta \leftarrow 0$
- ▶ Loop for each $s \in S$:
 - ▶ $v \leftarrow V(s)$
 - ▶ $V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s))[r + \gamma V(s')]$
 - ▶ $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
- ▶ until $\Delta < \theta$

S3: Policy Improvement

- ▶ $policy-stable \leftarrow 1$
- ▶ For each $s \in S$:
 - ▶ $old-action \leftarrow \pi(s)$
 - ▶ $\pi(s) \leftarrow \arg \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$
 - ▶ If $old-action \neq \pi(s)$, then $policy-stable \leftarrow 0$
- ▶ If $policy-stable$, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Exercise: Policy iteration for action values

Revise the algorithm on the previous page to use action values q instead and thus find q_* .

Value iteration

- ▶ A problem with policy iteration is the need to perform policy evaluation to some accuracy before improving the policy.
- ▶ Value iteration skips finding the policy and just updates the value to the maximum value

$$v(s) \leftarrow \max_a \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s, A_t = a]$$

Value iteration

- ▶ A problem with policy iteration is the need to perform policy evaluation to some accuracy before improving the policy.
- ▶ Value iteration skips finding the policy and just updates the value to the maximum value

$$v(s) \leftarrow \max_a \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s, A_t = a]$$

- ▶ The algorithm terminates when the size of the updates to any $v(s)$ is smaller than some threshold

Phew!

